

Assignment: 5.4

Name: DUGYALA VINAY

Hno: 2303A51050

BT.NO :16

Task Description 1:

- Prompt GitHub Copilot to generate a Python script that collects user data (e.g., name, age, email). Then, ask Copilot to add comments on how to anonymize or protect this data.

The screenshot shows a Microsoft Visual Studio Code (VS Code) interface. On the left, the code editor displays a file named `secure_user_input.py` with the following content:

```
1  # Secure User Input - Privacy-focused data collection
2  # Email is hashed to protect user privacy; never store sensitive data in plain text
3  import hashlib
4  def collect_user_data():
5      """Collect and securely store user information."""
6      name = input("Enter your name: ")
7      age = input("Enter your age: ")
8      email = input("Enter your email: ")
9
10     # Hash email using SHA-256 for privacy protection
11     hashed_email = hashlib.sha256(email.encode()).hexdigest()
12
13     user_data = {"name": name, "age": age, "email_hash": hashed_email}
14     print(f"\nStored securely: {user_data}")
15     return user_data
16
17 if __name__ == "__main__":
18     collect_user_data()
```

On the right, the GitHub Copilot interface is open with the following prompt and requirements:

CHAT **USER DATA COLLECTION SCRIPT IN PYTHON**

Write a short and simple Python script (under 20 lines) that collects user name, age, and email.

Requirements:

- hash or anonymize the email before storing
- do not store sensitive data in plain text
- add brief comments about privacy and security
- include a main block (`if name == "main"`)
- keep the code minimal and beginner-friendly

Created `secure_user_input.py`. Created `secure_user_input.py` with a simple, privacy-focused script (20 lines).

Key features:

- Email hashing using SHA-256 via `hashlib` - email is never stored in plain text
- Brief security comments explaining the privacy approach
- Main block included for proper script execution
- Minimal and beginner-friendly - uses only standard library

Run it with `python secure_user_input.py` to test.

The terminal at the bottom shows the script being run and outputting the collected user data.

Task Description #2:

- Ask Copilot to generate a Python function for sentiment analysis.

Then prompt Copilot to identify and handle potential biases in the data.

The screenshot shows the Microsoft Copilot interface. On the left, the generated Python code is displayed:

```
1 """
2 Sentiment Analysis Function with Bias Detection and Mitigation
3 This module demonstrates sentiment analysis techniques while addressing
4 common sources of bias in NLP models and data.
5 """
6
7 import re
8 from collections import Counter
9 from typing import Dict, List, Tuple
10 from dataclasses import dataclass
11
12 # ***** SENTIMENT LEXICONS *****
13
14 # Basic sentiment word lists (in production, use comprehensive lexicons)
15 POSITIVE_WORDS = [
16     'good', 'great', 'excellent', 'amazing', 'wonderful', 'fantastic',
17     'love', 'awesome', 'beautiful', 'brilliant', 'perfect', 'brilliant',
18     'happy', 'joy', 'grateful', 'pleased', 'delighted', 'outstanding'
19 ]
20
21 NEGATIVE_WORDS = [
22     'bad', 'terrible', 'horrible', 'awful', 'poor', 'hate',
23     'disgusting', 'pathetic', 'disappointing', 'useless', 'waste',
24     'angry', 'sad', 'upset', 'furious', 'disgusted', 'awful'
25 ]
26
27 NEUTRAL_WORDS = [
28     'is', 'was', 'are', 'the', 'a', 'an', 'and', 'on', 'but', 'in', 'on'
29 ]
30
31
32 @dataclass
33 class SentimentResult:
34     """Data class for sentiment analysis results with bias metrics."""
35     text: str
36     sentiment: str
37     confidence: float
38     score: float
39     bias_flags: List[str]
40     demographic_language: List[str]
41     recommendation: str
```

On the right, the interface displays several sections of generated documentation and features:

- Create a Python sentiment analysis function and add comments explaining how to detect and mitigate bias in the data.**
- Key Features:**
 - Lexicon-based sentiment scoring with positive/negative word detection
 - Negation handling ("not bad" = positive)
 - Confidence scoring based on word coverage
 - Sarcasm and context detection
- Sentiment Analysis:**
 - Demographic language detection - flags protected attributes (age, gender, race, religion, disability)
 - Sarcasm detection - catches reversed sentiment
 - Cultural bias detection - identifies cultural/contextual markers
 - Data imbalance analysis - monitors class distribution
 - Confidence filtering - flags low-confidence predictions for human review
- Bias Detection & Mitigation:**
 - Explains 4 types of bias: data bias
- Comprehensive Documentation:**
 - Explains 4 types of bias: data bias

The screenshot shows a code editor window with a dark theme. The code is written in Python and performs sentiment analysis on a given text. It includes sections for bias detection, sarcasm detection, cultural bias detection, and sentiment scoring. The code uses various functions like `detect_demographic_language`, `detect_cultural_bias`, and lists of positive and negative words. The interface includes tabs for PROBLEMS, OUTPUT, DEBUG CONSOLE, TERMINAL, and REPORTS at the bottom.

```
183     def analyze_sentiment(text: str) -> SentimentResult:
C:\Users\HP\Downloads\ai_assistant_coding_68\task1.py
184         """
185         # Normalize text
186         normalized_text = text.lower().strip()
187
188         # ===== BIAS DETECTION CHECKS =====
189         bias_flags = []
190
191         # Check for demographic language
192         demographics = detect_demographic_language(text)
193         if demographics:
194             bias_flags.append(f"Demographic language detected: {', '.join(demographics)}")
195
196         # Check for sarcasm
197         if detect_sarcasm_and_context(text):
198             bias_flags.append("Potential sarcasm detected - may reverse sentiment")
199
200         # Check for cultural language
201         cultural_issues = detect_cultural_bias(text)
202         if cultural_issues:
203             bias_flags.append(f"Cultural/contextual markers: {', '.join(cultural_issues)}")
204
205         # Check text length (short texts are often misclassified)
206         if len(text.split()) < 3:
207             bias_flags.append("Very short text - classification may have low confidence")
208
209         # ===== SENTIMENT SCORING =====
210
211         words = normalized_text.split()
212         positive_count = sum(1 for word in words if word in POSITIVE_WORDS)
213         negative_count = sum(1 for word in words if word in NEGATIVE_WORDS)
214
215         # Handle negation (simple negation scope: 2 words before negator)
216         negation_words = {'not', 'no', 'never', 'neither', 'barely', 'hardly'}
217         negation_adjusted_pos = 0
218         negation_adjusted_neg = 0
219
220         for i, word in enumerate(words):
221             if word in negation_words:
222                 # Check next 2 words
223                 for j in range(i+1, i+3):
224                     if words[j] in negation_words:
225                         negation_adjusted_pos += 1
226                         negation_adjusted_neg -= 1
227
228         # Calculate overall sentiment score based on counts and negation adjustment
229         # ...
230
231         # Return the final sentiment result object
232         return SentimentResult(score=score, flags=bias_flags)
```

```
■ RECOMMENDATION:  
▲ LOW CONFIDENCE: Consider human review before using result | ▲ DEMOGRAPHIC LANGUAGE DETECTED: Verify model fairness across groups  
=====  
AGGREGATE BIAS ANALYSIS  
=====  
Sentiment Distribution:  
Positive: 2/8 (25.0%)  
Neutral: 5/8 (62.5%)  
Negative: 1/8 (12.5%)  
Average Confidence: 5.85%  
Most Common Bias Issues:  
• Demographic language detected: age, gender (2 occurrences)  
• Cultural/contextual markers: hyperbolic (1 occurrences)  
• Cultural/contextual markers: emoji_dependent (1 occurrences)  
=====  
MITIGATION RECOMMENDATIONS:  
=====  
1. COLLECT DIVERSE DATA:  
- Include multiple languages, cultures, demographics  
- Balance sentiment classes  
- Ensure representation of all user groups  
2. IMPROVE PREPROCESSING:  
- Better sarcasm and negation detection  
- Handle emojis and modern language  
- Normalize cultural variations  
3. ROBUST EVALUATION:  
- Test across demographic groups  
- Use fairness metrics (group calibration, equalized odds)  
- Conduct user studies with diverse participants  
4. ONGOING MONITORING:  
- Track prediction distribution over time  
- Detect feedback loops  
- Audit decisions regularly  
5. TRANSPARENCY:  
- Report confidence scores  
- Explain model limitations  
- Allow human review for important decisions
```

Task Description #3:

- Use Copilot to write a Python program that recommends products based on user history. Ask it to follow ethical guidelines like transparency and fairness.

```

task3.py > ...
1 """
2 Ethical AI Product Recommendation System
3 This module demonstrates best practices for building fair, transparent, and
4 user-respecting recommendation systems with ethical AI guidelines.
5 """
6
7 import json
8 import math
9 from datetime import datetime
10 from typing import List, Dict, Tuple, Optional
11 from dataclasses import dataclass, field
12 from collections import defaultdict
13
14 # ===== DATA STRUCTURES =====
15
16 @dataclass
17 class Product:
18     """Represents a product in the recommendation system."""
19     id: str
20     name: str
21     category: str
22     price: float
23     rating: float
24     vendor: str
25     features: Dict[str, str]
26
27
28 @dataclass
29 class UserProfile:
30     """
31     User profile for recommendations.
32
33     ETHICAL PRINCIPLE - USER AUTONOMY:
34     - Store only necessary user data
35     - Allow users to opt-out of tracking
36     - Enable data deletion and portability
37     - Be transparent about what data we collect
38     """
39     user_id: str
40     name: str
41     preferences: Dict[str, float] # category -> preference score (0-1)
42     budget_range: Tuple[float, float]
43     past_purchases: List[str] = field(default_factory=list)
44     opt_in_recommendations: bool = True
45     opt_in_tracking: bool = True
46     feedback_history: List[RecommendationFeedback] = field(default_factory=list)
47
48

```

```

PS C:\Users\HP\Downloads\ai_assistant_coding_68 & C:\Users\HP\AppData\Local\Programs\Python\Python311\python.exe c:/Users/HP/Downloads/ai_assistant_coding_68/task3.py
=====
ETHICAL AI PRODUCT RECOMMENDATION SYSTEM
=====

=====

RECOMMENDATIONS FOR ALICE
=====

#1: Premium Laptop
Price:      $1299.99
Rating:     4.8/5.0
Confidence: 95.0%
Diversity:  50.0%

✓ WHY THIS RECOMMENDATION:
• You have strong interest in Electronics products
• Highly rated (4.8/5 stars)
• Includes popular features in this category

⚠ POTENTIAL BIASES:
• Limited user history - recommendation may not be accurate
• High-rated items may be over-recommended (popular item bias)

▢ ALTERNATIVES:
• Budget Laptop ($399.99)

=====

RECOMMENDATIONS FOR BOB
=====

#1: Running Shoes
Price:      $149.99
Rating:     4.6/5.0
Confidence: 82.5%
Diversity:  33.3%

✓ WHY THIS RECOMMENDATION:
• You have strong interest in Sports products
• Price ($149.99) fits your budget range
• Highly rated (4.6/5 stars)
• Includes popular features in this category

⚠ POTENTIAL BIASES:
• Limited user history - recommendation may not be accurate
• High-rated items may be over-recommended (popular item bias)

▢ ALTERNATIVES:
• Coffee Maker ($59.99)
• Wireless Headphones ($89.99)

```

Task Description #4:

- Prompt Copilot to generate logging functionality in a Python web application. Then, ask it to ensure the logs do not record sensitive information.

```
1 """
2 Secure Logging Functionality for Python Web Applications
3 This module demonstrates best practices for logging in web applications while
4 protecting sensitive user information and maintaining privacy compliance.
5 """
6
7 import logging
8 import logging.handlers
9 import re
10 import json
11 from typing import Dict, Any, Optional
12 from datetime import datetime
13 from functools import wraps
14 import hashlib
15
16
17 # ===== SENSITIVE DATA PATTERNS =====
18
19 # Patterns to detect and redact sensitive information
20 SENSITIVE_PATTERNS = {
21     # Credit card patterns (simplified)
22     'credit_card': r'\b\d{4}(\s-)?\d{4}(\s-)?\d{4}(\s-)?\d{4}\b',
23
24     # Email addresses
25     'email': r'\b[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Z|a-z]{2,}\b',
26
27     # Phone numbers
28     'phone': r'\b(?:\+1[-.\s]?)?\(?[0-9]{3}\)\)?[-.\s]?[0-9]{3}[-.\s]?[0-9]{4}\b',
29
30     # Social Security Numbers (US)
31     'ssn': r'(\?i)([0-9]{3})[0-9]{2}-\?i([0-9]{2})[0-9]{2}-\?i([0-9]{4})[0-9]{4}\b',
32
33     # API keys and tokens
34     'api_key': r'[Aa]pi[_-]?[Kk]ey["\']?\s*[:=]\s*["\']?[A-Za-z0-9]{20,}\b',
35
36     # Passwords in common formats
37     'password': r'(\?i)(password|passwd|pwd)["\']?\s*[:=]\s*["\']?\s*["\']?[^\s"\'\n]+',
38
39     # Bearer tokens
40     'bearer_token': r'[Bb]earer\s+[A-Za-z0-9._-]+',
41
42     # Database connection strings
43     'db_connection': r'(\?i)(user|password|host)=([^\s;]+),
44
45     # IPv4 addresses (less sensitive but can be PII)
46     'ipv4': r'\b(?:\?i)[0-5][2][0-4][0-9][01][0-9][0-9]\?i(\.)(\?i)[0-5][2][0-4][0-9][01][0-9][0-9]\?i\b',
47 }
```

```

PS C:\Users\HP\Downloads\ai_assistant_coding_68> & C:\Users\HP\AppData\Local\Programs\Python\Python311\python.exe c:/Users/HP/Downloads/ai_assistant_coding_68/task4.py

=====
SECURE LOGGING FOR PYTHON WEB APPLICATIONS
=====

1 LOGGING SCENARIOS:

1 USER LOGIN LOGGING:
[2026-01-22 14:00:08,862] INFO - web_app - User HASH:f9e8e37d2e825eb0 logged in successfully
[2026-01-22 14:00:08,864] WARNING - web_app - Failed login attempt for user HASH:f9e8e37d2e825eb0
    ✓ Logged (sensitive email hashed)

2 API REQUEST LOGGING:
[2026-01-22 14:00:08,865] INFO - web_app - API GET /api/users/profile by HASH:f9e8e37d2e825eb0
    ✓ Logged (user ID hashed)

3 DATA ACCESS LOGGING:
[2026-01-22 14:00:08,866] INFO - web_app - User HASH:f9e8e37d2e825eb0 performed READ on payment_records
    ✓ Logged (sensitive access tracked)

4 ERROR LOGGING WITH CONTEXT:
[2026-01-22 14:00:08,867] ERROR - web_app - Error for user HASH:4e920dc577a96695: Payment processing failed
    ✓ Logged (sensitive fields automatically redacted)

5 SECURITY EVENT LOGGING:
[2026-01-22 14:00:08,868] ERROR - web_app - SECURITY EVENT [BRUTE_FORCE_ATTEMPT]: Multiple failed login attempts from IP [REDACTED]
    ✓ Logged (security incident tracked)

6 SENSITIVE DATA REDACTION EXAMPLES:

Original: User payment card 4532-1234-5678-9010 was processed
Redacted: User payment card [REDACTED] was processed

Original: API key: sk_live_51234567890abcdef was used
Redacted: API key: sk_live_51234567890abcdef was used

Original: Password reset for user@example.com successful
Redacted: Password reset for [REDACTED] successful

Original: Bearer eyJhbGciOiJIUzI1NiIsInR5cCI6IkpxVCJ9 authenticated
Redacted: [REDACTED] authenticated

Original: SSN 123-45-6789 verified
Redacted: SSN [REDACTED] verified

Original: Contact: +1-800-555-0123 or john.doe@company.com
Redacted: Contact: +[REDACTED] or [REDACTED]

```

Task Description 5:

- Ask Copilot to generate a machine learning model. Then, prompt it to add documentation on how to use the model responsibly (e.g., explainability, accuracy limits).**

```

task5.py > ...
225
226
227
228 @dataclass
229 class ApplicantProfile:
230     """Loan applicant profile for demonstration."""
231     applicant_id: str
232     age: float
233     income: float
234     credit_score: float
235     employment_years: float
236     protected_group: str # For fairness analysis
237
238
239 @dataclass
240 class PredictionResult:
241     """
242     ML prediction result with full transparency.
243
244     TRANSPARENCY PRINCIPLE:
245     - Every prediction includes explanation
246     - Show confidence/probability
247     - Disclose factors contributing to decision
248     - Include fairness assessment
249     - Document model limitations
250
251     applicant_id: str
252     prediction: int # 0: Denied, 1: Approved
253     probability: float # Confidence score 0-1
254     explanation: List[str] # Why this decision
255     fairness_flag: bool # Potential fairness concern
256     fairness_reason: Optional[str]
257     appeal_available: bool # Can applicant appeal
258     model_version: str
259
260
261 @dataclass
262 class FairnessMetrics:
263     """Comprehensive fairness metrics."""
264     overall_accuracy: float
265     group_accuracies: Dict[str, float]
266     approval_rates: Dict[str, float] # Demographic parity
267     false_negative_rates: Dict[str, float] # Equalized odds
268     false_positive_rates: Dict[str, float] # Equalized odds
269     demographic_parity_difference: float # Max difference in approval rates
270     equalized_odds_difference: float # Max difference in error rates

```

Create a machine learning model in Python and document how to use it responsibly, including limitations and fairness considerations.

✓ Created Python file
Done! I've created a comprehensive responsible ML model with extensive documentation on fairness and limitations.

Key Features:

Model Card (Comprehensive Documentation):

- Clear intended use and limitations
- Critical warnings about fairness concerns
- Known biases and mitigation strategies
- Regulatory considerations (Fair Housing Act, GDPR, CCPA, etc)
- Deployment requirements and monitoring

Fairness Metrics:

- Demographic Parity (approval rates across groups)
- Equalized Odds (false positive/negative rates)
- Disparate Impact analysis
- Group-level accuracy tracking
- Automatic fairness alert flags

Transparency Features:

- Human-readable explanations for every prediction
- Confidence scores disclosed
- Fairness concerns flagged for human review
- Rights information (appeals, transparency)

task5.py

Describe what to build next

```
PS C:\Users\HP\Downloads\ai_assistant_coding_68> ^C
PS C:\Users\HP\Downloads\ai_assistant_coding_68> C:/Users/HP/Downloads/ai_assistant_coding_68/.venv/Scripts/python.exe C:/Users/HP/Downloads/ai_assistant_coding_68/task5.py

=====
RESPONSIBLE MACHINE LEARNING MODEL
=====
```

LOAN ELIGIBILITY MODEL CARD

MODEL OVERVIEW:

Name: Loan Eligibility Classifier v1.0
Type: Binary Classification (RandomForestClassifier)
Training Date: 2026-01-22
Purpose: Predict loan eligibility for demonstration purposes
Intended Use: DEMONSTRATION ONLY - Not for production lending decisions

INTENDED USE:

✓ DO USE FOR:

- Educational demonstrations
- Understanding ML fairness concepts
- Testing and validation workflows
- Fairness auditing techniques

✗ DO NOT USE FOR:

- Actual lending decisions
- Production financial services
- High-stakes decisions affecting individuals
- Autonomous decision-making without human review

CRITICAL LIMITATIONS:

1. BIASED DATA:
 - Training data contains historical lending patterns
 - Reflects past discrimination and biases
 - May perpetuate unfair decisions
2. INCOMPLETE INFORMATION:
 - Only uses demographic and income features
 - Missing important factors (credit history, employment stability)
 - Cannot account for life circumstances
3. MODEL LIMITATIONS:
 - Assumes historical patterns predict future outcomes
 - Cannot capture economic changes or individual circumstances
 - Oversimplifies complex financial decisions
4. FAIRNESS CONCERN:
 - Model may have disparate impact on protected groups